


Recall: Deployment is closing the loop

Programs that are never deployed have not fulfilled their purpose. We must deploy!

To do so we must answer:

- Is our application in a working state?
- Do we have the necessary HW/SW resources?
- How do we actually deploy?

Recall: CI, CD and more



CI rigorously tests every integration in production-like environment

- Prevent development-production mismatch
- Test multiple browsers, etc.
- “Stress test” code for performance, fault-tolerance, etc.

Then we deploy!

By deploying frequently, we make what was rare and fraught common and unremarkable!

Recall: DevOps principles

- Involve operations in each phase of a system's design and development,
- Heavy reliance on automation versus human effort,
- The application of engineering practices and tools to operations tasks

*aaS: _____ as code

Platform-as-a-
Service

*Three-tier
architecture as
code*

1. Deploy (that's it!)

Infrastructure-
as-a-Service

*"Infrastructure
as code"*

1. Configure (with tools like Ansible, etc.)
2. Deploy

Bare Metal

*Just
infrastructure*

1. Rack
2. Configure
3. Deploy

The *aaS division of labor

PaaS handles...	You handle...
“Easy” tiers of horizontal scaling	Minimize load on database
Component-level performance tuning	Application-level performance tuning (e.g., caching)
Infrastructure-level security	Application-level security

What about upgrades? Automation and rigorous processes in action

- Can't or don't want to rollout new feature simultaneously to all servers
 - Version n and $n+1$ will co-exist
- Naïve solution: Downtime
- Alternative: **Feature flags**
 1. Do non-destructive migration
 2. Deploy code protected by feature flag
 3. Flip feature flag on; if disaster, flip it back
 4. Once all records moved, deploy entirely new code
 5. Apply migration to remove old columns
- Other FF uses: A/B testing, ...

Kinds of monitoring

“If you haven't ~~tried~~ monitored it, assume it's broken.*”

- At development time (*profiling*)
Identify possible performance/stability problems *before* they get to production
- In production
Internal: Instrumentation embedded in application and/or framework
External: Active probing by other site(s)/tools.

Performance and security metrics

Availability or Uptime

What % of time is site up and accessible?

Responsiveness

How long after a click does user get response?

Scalability

As number users increases, can you maintain responsiveness without increasing cost/user?

Authorization (Privacy)

Is data access limited to the appropriate users?

Authentication

Can we trust that user is who s/he claims to be?

Data integrity

Is users' sensitive data tamper-evident?

Performance &
Stability

Security

Google's 4 "golden" signals

- Latency Can be confounded by errors. How?
Time to service a request
- Traffic Application specific metric: requests/s, I/O rate, ...
How much demand is being place on your system
- Errors
Rate of requests that fail
- Saturation
How "full" your system is (when will you hit ceiling?)

“Premature optimization is the root of all evil”*

- Users expect speed!
99 percentile matters, not just “average”
- There are lots of reasons for “too slow”
- Don’t assume, measure!
Monitoring is your friend: measure twice, cut once!

*Variously attributed to Hoare, Knuth, Dijkstra,

Simplified (& false) view of response time


For *normal distribution* of response times:

± 2 standard deviations around mean is 95% CI

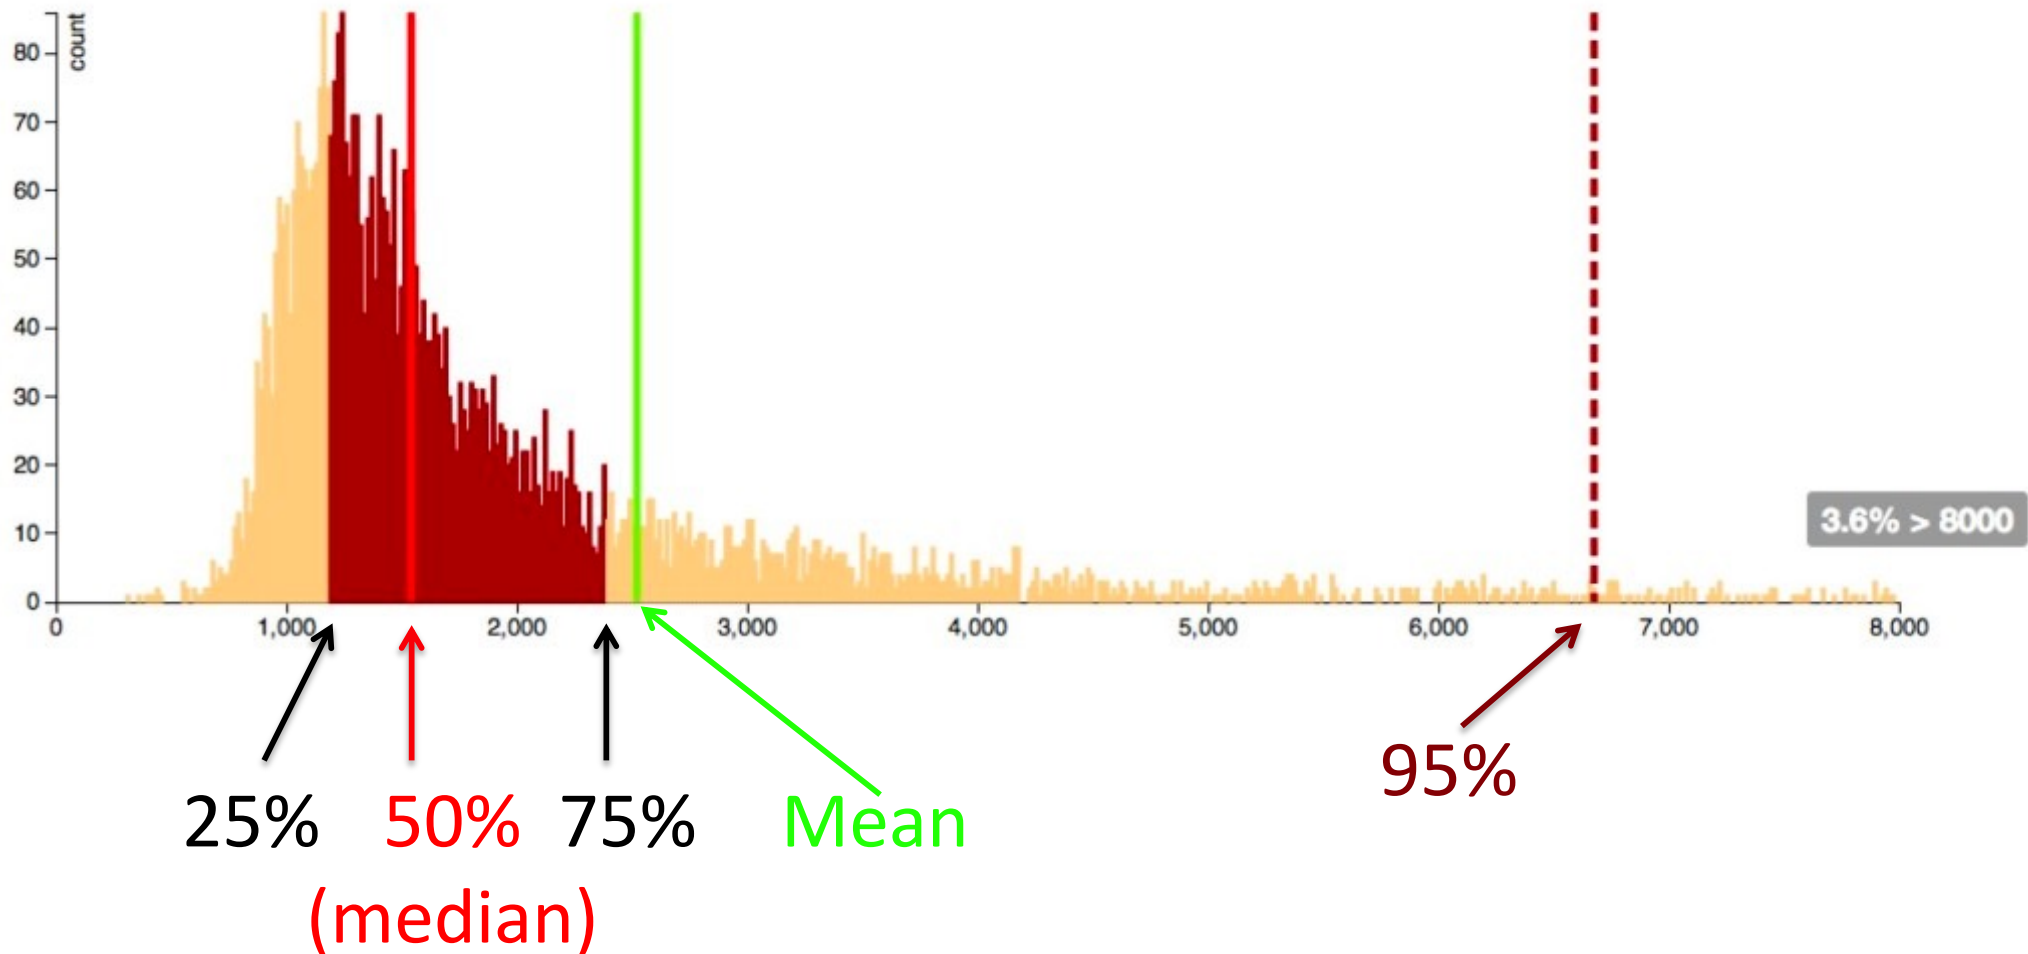
Average response time T
means:

95%ile users are getting $T+2\sigma$

99.7%ile users get $T+3\sigma$



A real example: The long tail



Service Level Objective (SLO): Target value for your service

Instead of worst case or average metric, specify a percentile, target and window

99% of requests complete in < 1 second, averaged over a 5 min. window

SLOs set customer expectations

Make sure you have a safety margin

Overachieving can be problematic too! **How?**

Service Level Agreements (SLAs) attach contractual obligations to SLOs

How can you fix “slow”?

- Add more resources, i.e., over-provision
 - Easy to scale presentation and logic tiers for small sites (readily automated in the “cloud”)
 - More expensive for larger sites (10% of 10,000 machines is a big number!)
- Make your application more efficient
 - Most effective when there is one bottleneck

The fastest computation is the one you don't do

- Don't forget big-O and CS fundamentals, e.g.

`Array.include` VS. `Set` for unique

Smart use of DB indexes

- Caching (and memoization more generally)
- Avoid “toxic” queries, e.g.
 - “n+1” query for associations

DB is one of the hardest components to scale, aim to *be kind to your database*.

Indexes: $O(< n)$ queries

Index is a tree, hash-table or other data structure optimized for efficient queries

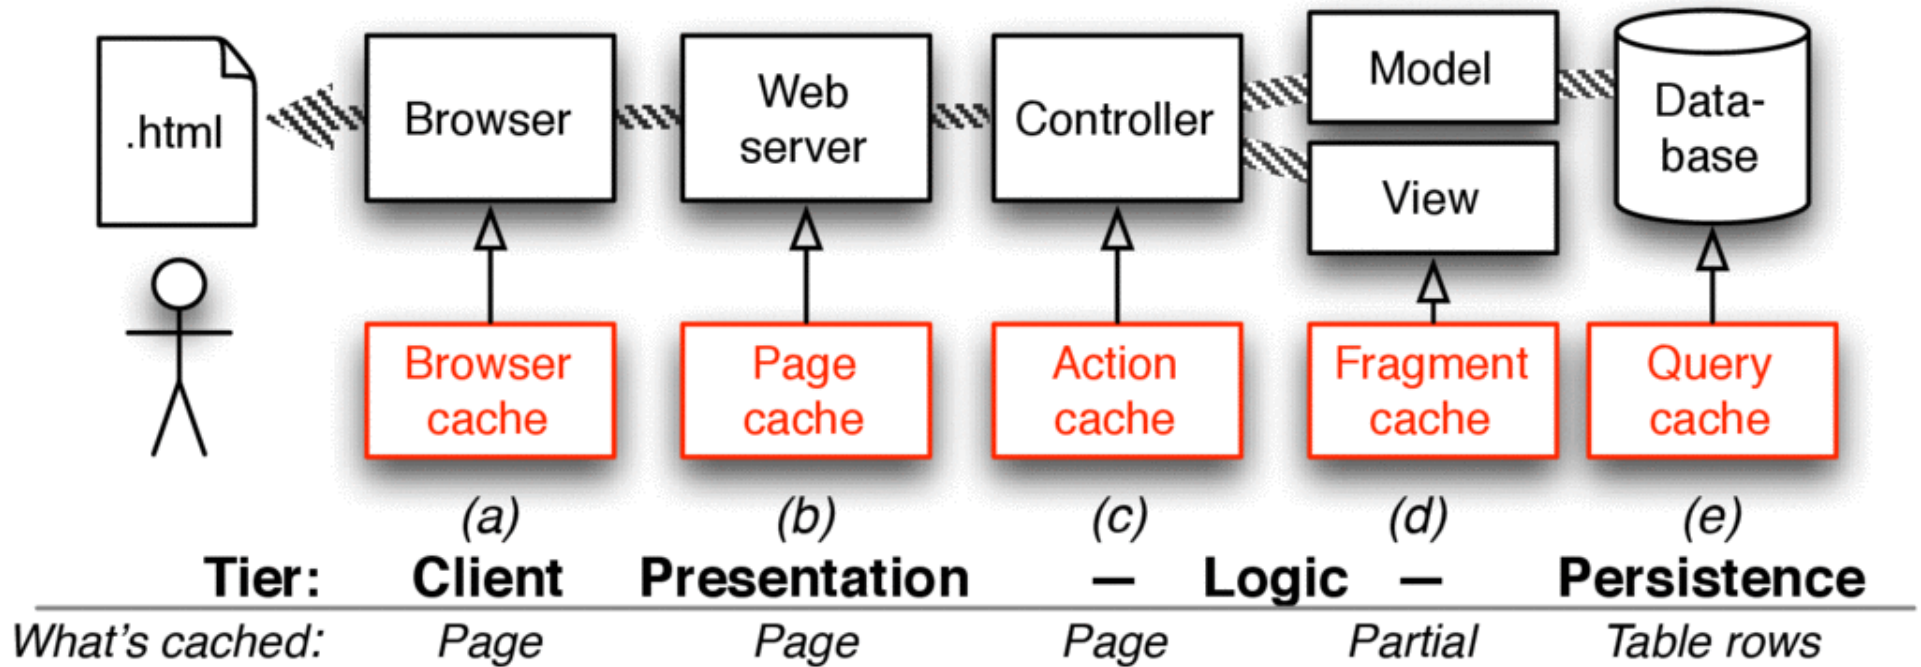
# of reviews:	2000	20,000	200,000
Read 100, no indices	0.94	1.33	5.28
Read 100, FK indices	0.57	0.63	0.65
Performance	166%	212%	808%

Sub-linear scaling!

Why not use an index for every field?

- Requires additional storage space for each index
- Slows down insert/edit (need to update the index)

Cache what hasn't changed



“There are only two hard things in Computer Science: cache invalidation and naming things.” –Phil Karlton

n+1 queries (or leaky abstractions)

Recall in the Film Explorer a user “has many” films “through” ratings

```
User.query().where('zip', '05753').then((fans) => {  
  fans.forEach((fan) => {  
    fan.$relatedQuery('films')...  
  });  
});
```

1 query for each user (i.e. $n+1$ queries for n users)

More subtle for other ORMs, e.g. `fan.films()` is really a query

```
User.query()  
  .where('zip', '05753')  
  .eager('films')  
  .then((fans) => {  
    fans.forEach((fan) => {  
      fan.films ...  
    });  
  });
```

Just 1 or 2 queries, but DB “leaking” through ORM abstraction